

Advanced SQL

02 — The Core of SQL

Summer 2022

Torsten Grust
Universität Tübingen, Germany

1 | The Core of SQL

- Let us recollect the **core constructs of SQL**, synchronize notation, and introduce query conventions.
- If you need to refresh your SQL memory, consider
 - the notes for [DB1](#) (Chapters 6 and 9)
 - the [PostgreSQL web](#) (Part II, The SQL Language)
- We will significantly expand on this base SQL vocabulary during the semester.

Sample Table

Table **T** serves as a common “playground” for the upcoming SQL queries:

<u>a</u>	<u>b</u>	<u>c</u>	<u>d</u>
1	'x'	true	10
2	'y'	true	40
3	'x'	false	30
4	'y'	false	20
5	'x'	true	NULL

Table **T**

```
CREATE TABLE T (a int PRIMARY KEY, -- implies NOT NULL
                  b text,           -- here: char(1)
                  c boolean,
                  d int);
```

2 : Row Variables

- Iterate over all rows of table **T** (in *some* order: bag semantics), bind **row variable** **t** to current row:

```
SELECT t          -- 2 t is bound to current row of T
FROM   T AS t     -- 1 bind/introduce t
```

- If you omit **AS t** in the **FROM** clause, a row variable **T** (generally: **AS <table name>**) will be implicitly introduced.
- This course: always explicitly introduce/name row variables for disambiguation, clarity, readability.

Row Values

```
SELECT t      -- 2 t is bound to current row of T
FROM   T AS t -- 1 bind/introduce t
```

- Row variable `t` is iteratively bound to **row values** whose field values and types are determined by the rows of table `T`:

field names:

	a	b	c	d	
	↓	↓	↓	↓	
t ≡	(5,	'x',	true,	NULL)	} row values
t ≡	(1,	'x',	true,	10)	
⋮					
t ≡	(2,	'y',	true,	40)	
	↑	↑	↑	↑	

field types:

int text boolean int

Row Types

- `t :: T` with `T = (a int, b text, c boolean, d int)`.¹ **Row type** `T` is defined when `CREATE TABLE T (...)` is performed.
- A row type `τ` can also be explicitly defined via

```
CREATE TYPE τ AS (a int, b text, c boolean, d int)
```

- A table `T1` equivalent to `T` — well, almost... — may then be created via

```
CREATE TABLE T1 OF τ
```

¹ Read `::` as “has type.”

Row Field Access and * (“Star”)

- Named **field access** uses dot notation. Assume `t :: T` and binding `t ≡ (5, 'x', true, NULL)` then:
 - `t.b` evaluates to `'x'` (of type `text`),
 - `t.d` evaluates to `NULL` (of type `int`).
- Field names are *not* first-class in SQL and must be provided verbatim (i.e., may *not* be computed).
- Notation `t.*` abbreviates `t.a, t.b, t.c, t.d` in contexts where this makes sense.²

² `t.*` is most often used in `SELECT` clauses.

Row Comparisons

- **Row comparisons** between rows t_1 , t_2 are performed field-by-field and lexicographically (provided that the field types match). Assume $t_1 :: T$, $t_2 :: T$:
 - $t_1 = t_2 \iff t_1.a = t_2.a \text{ AND } \dots \text{ AND } t_1.d = t_2.d$
 - $t_1 < t_2 \iff$
 $t_1.a < t_2.a \text{ OR } (t_1.a = t_2.a \text{ AND } t_1.b < t_2.b) \text{ OR } \dots$
- A row value is **NULL** iff *all* of its field values are **NULL**.

Assume the binding $t \equiv (\text{NULL}, \text{NULL}, \text{NULL}, \text{NULL})$. Then $t \text{ IS NULL}$ holds.

3 | The **SELECT** Clause

A **SELECT** clause evaluates n expressions e_1, \dots, e_n :

```
SELECT  $e_1$  AS  $c_1$ , ...,  $e_n$  AS  $c_n$ 
```

- Creates n columns named c_1, \dots, c_n .
- In absence of **AS** c_i , PostgreSQL assigns name "**?column?**" (for *all* such unnamed columns) \Rightarrow ambiguity ☹.
- This course: explicitly use **AS** to name columns unless a name can be derived from e_i (e.g., when $e_i \equiv t.a$).
- If column/table names are case-sensitive or contain whitespace/symbols/keywords: wrap in quotes " c_i ".

Standalone **SELECT**

- If query Q generates n row bindings, **SELECT** is evaluated n times to emit n rows (but see *aggregates* below).
- A standalone **SELECT** (no **FROM** clause) is evaluated exactly once and emits a single row:

```
SELECT 1 + 41 AS "The Answer", 'Gla' || 'DOS' AS Portal;
```

The Answer	portal
42	GlaDOS

4 : Literal Tables (**VALUES**)

A **VALUES** clause constructs a transient table from a list of provided **row values** e_i :

VALUES e_1, \dots, e_n

- If $n > 1$, the e_i must agree in arity and field types (row value e_1 is used to infer and determine types).
- **VALUES** automatically assigns column names "column< i >". Use column aliasing to assign names (see **FROM** below).
- **Orthogonality:** a **VALUES** clause (in parentheses (...)) may be used anywhere a SQL query expects a table.

5 : Generating Row Variable Bindings (**FROM**)

A **FROM** clause expects a set of tables T_i and successively binds the row variables t_i to the tables' rows:

```
SELECT ...           -- 2
FROM    $T_1$  AS  $t_1$ , ...,  $T_n$  AS  $t_n$   -- 1
```

- The T_i may be table names or SQL queries computing tables (enclosed in (...)).
- If you need to rename the columns of T_i (recall the **VALUES** clause), use **column aliasing** on all (or only the first k ☹) columns:

T_i AS $t_i(\underline{c_{i,1}}, \dots, \underline{c_{i,k}})$

FROM Computes Cartesian Products

```
SELECT ...  
FROM     $T_1$  AS  $t_1$ , ...,  $T_n$  AS  $t_n$ 
```

- This **FROM** clause generates $|T_1| \times \dots \times |T_n|$ bindings.
Semantics: compute the **Cartesian product** $T_1 \times \dots \times T_n$, draw the bindings for the t_i from this product.
- **FROM** operates over a *set* of tables ('**,**' is associative and commutative).
- In particular, row variable t_i is *not* in scope in the table subqueries T_{i+1}, \dots, T_n (but see **LATERAL**).

6 : WHERE Discards Row Bindings

A **WHERE** clause introduces a predicate p that is evaluated under all row variable bindings generated by **FROM**:

SELECT	...	--	3
FROM	T_1 AS t_1 , ..., T_n AS t_n	--	1
WHERE	p	--	2

- All row variables t_i are in scope in p .
- Only bindings that yield $p = \text{true}$ are passed on.³
- Absence of a **WHERE** clause is interpreted as **WHERE true**.

³ If p evaluates to **NULL** \neq **true**, the binding is discarded.

7 : Compositionality: Subqueries Instead of Values

“The meaning of a complex expression is determined by the meanings of constituent expressions.”

—Principle of Compositionality

With the advent of the SQL-92 and SQL:1999 standards, SQL has gained in **compositionality** and **orthogonality**:


- Wherever a (tabular or scalar) value *v* is required, a SQL expression in *(...)*—a **subquery**—may be used to compute *v*.
- Subqueries nest to arbitrary depth.

Scalar Subqueries: Atomic Values

A SQL query that computes a **single-row, single-column table** (column name \square irrelevant) may be **used in place of an atomic value** v :



In a **scalar subquery**...

- ... an empty table is interpreted as **NULL**,
- ... a table with > 1 rows or > 1 columns will yield a **runtime error** .

Scalar Subqueries: Atomic Values

generate single column

↓

```
SELECT 2 + (SELECT t.d AS _
             FROM T AS t
             WHERE t.a = 2) AS "The Answer"
```

equality predicate on key column,
will yield ≤ 1 rows

- **Runtime errors** for `WHERE t.a > 2` or `SELECT t.a, t.d`
- Subquery yields `NULL`: `WHERE t.a = 0`
- `AS _` assigns a “*don't care*” column name—this is a case where column naming is obsolete and adds nothing.

Scalar Subqueries: Row Values

A SQL query that computes a **single-row table** with column names c_i may be **used in place of row value** (v_1, \dots, v_n) with field names c_i :

c_1	\dots	c_n
v_1	\dots	v_n


In a **scalar subquery**...

- ... an empty table is interpreted as $(\text{NULL}, \dots, \text{NULL})$,
- ... a table with > 1 rows will yield a **runtime error**.


Row Variable Scoping

Subqueries may **refer to any row variable t bound in their enclosing queries** (up to the top-level query):

enclosing query



enclosing query



- NB: Seen from inside the subquery—*i.e.*, inside the (...)—row variable t is *free* (binding provided externally).

Subqueries, Free Row Variables, Correlation

- If t is free in subquery q , we may understand the subquery as a function $q(t)$: you supply a value for t , I will compute the (tabular) value of q :

<pre> SELECT t1.* FROM T AS t1 WHERE t1.b <> (SELECT t2.b FROM T AS t2 WHERE t1.a = t2.a) </pre>	}	<p>evaluated 5 times under t1 bindings:</p> <p>t1 ≡ (1, ...)</p> <p>t1 ≡ (2, ...)</p> <p>t1 ≡ (3, ...)</p> <p>t1 ≡ (4, ...)</p> <p>t1 ≡ (5, ...)</p>
<p>↑</p> <p>free</p>		

- Subqueries featuring free variables are also known as **correlated**.

8 : Row Ordering (**ORDER BY**)

SQL tables are **unordered bags** of rows, but rows may be **locally ordered** for result display or positional access:

```
SELECT ...           -- 3
FROM   ...           -- 1
WHERE  ...           -- 2
ORDER BY  $e_1, \dots, e_n$  -- 4
```

- The order of the e_i matters: sort order is determined lexicographically with e_1 being the major criterion.
- Sort criteria e_i are expressions that may refer to column names in the **SELECT** clause (evaluated after **SELECT**).

SELECT t.* FROM T AS t ...

a	b	c	d
5	'x'	true	NULL
1	'x'	true	10
4	'y'	false	20
3	'x'	false	30
2	'y'	true	40

... ORDER BY t.d ASC NULLS FIRST

a	b	c	d
4	'y'	false	20
2	'y'	true	40
3	'x'	false	30
1	'x'	true	10
5	'x'	true	NULL

... ORDER BY t.b DESC, t.c

- Note: **ASC** (ascending) is default. **NULL** is larger than any non-**NULL** value. Ties: order is implementation-dependent.

Row Order is Local Only


ORDER BY establishes a well-defined row order that is **local** to the current (sub)query:

may yield rows in any order

↓

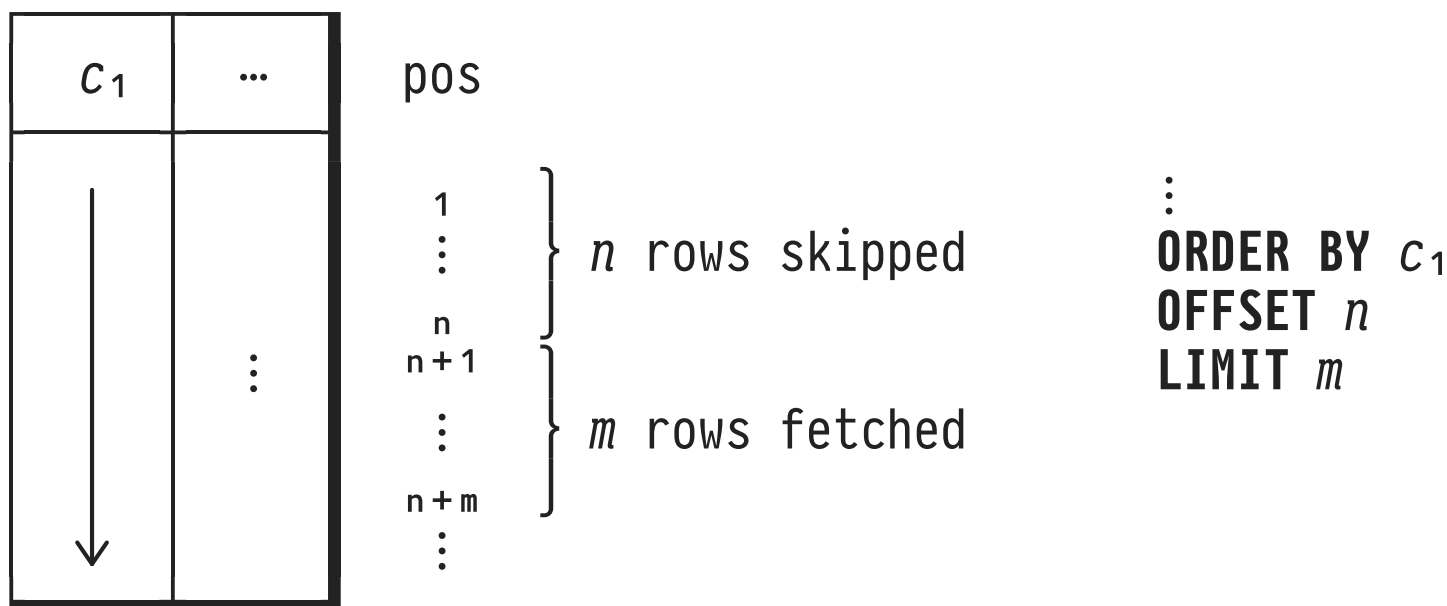
```
SELECT t1.*  
FROM   (SELECT t2.*  
        FROM   T AS t2  
        ORDER BY t2.a) AS t1;
```

} guaranteed **row order**
inside the subquery **only**

-  Never rely on row orders that appear consistent across runs—may vary between DBMSs, presence of indexes, etc.
- **Q:** What, then, is such local row order good for?

Positional Access to Rows

Once row order has been established it makes sense to “skip to the n^{th} row” or “fetch the **next** m rows.”




- **OFFSET 0**: read from the start. **LIMIT ALL**: fetch all rows.
- Alternative syntax: **FETCH [FIRST | NEXT] m ROWS ONLY**.

9 : Identify Particular Rows Among Peers (**DISTINCT ON**)

Extract the **first row** among a group of equivalent rows:

	prefix of ORDER BY clause	
SELECT DISTINCT ON	4 (e_1, \dots, e_n)	-- 2
FROM	\dots	-- 1
ORDER BY	$e_1, \dots, e_n, e_{n+1}, \dots, e_m$	-- 3

1. Sort rows in $e_1, \dots, e_n, e_{n+1}, \dots, e_m$ order.
 2. Rows with identical e_1, \dots, e_n values form one **group**.
 3. From each of these groups, pick **the first row** in e_{n+1}, \dots, e_m order.
-  Without **ORDER BY**, step 3 picks *any* row in each group.

DISTINCT ON: Group, Then Pick First in Each Group

SELECT DISTINCT ON (A_1) ... -- For each A_1 , pick the row ...
FROM ...
ORDER BY A_1 , A_2 **DESC** -- ... with the largest A_2

	A_1	A_2	...
	\vdots	\vdots	\vdots
group {	x_i	y_{i1}	...
	x_i	\vdots	\vdots
group {	x_j	y_{j1}	...
	x_j	\vdots	\vdots
	x_j	\vdots	\vdots
	\vdots	\vdots	\vdots

← pick
 } discard
 ← pick
 } discard

DISTINCT: Table-Wide Duplicate Removal

Keep only a single row from each group of **duplicates**:

```
SELECT DISTINCT ③  c1, ..., ck      -- ②
FROM      ...                -- ①
```

- True duplicate removal: rows are considered identical if they agree on **all** *k* columns *c_i*.⁴
- Row order is irrelevant. **DISTINCT** returns a *set of rows*.
- May use **SELECT ALL ...** to explicitly document that a query is expected to return duplicate rows.

⁴ This is equivalent to **SELECT DISTINCT ON (c₁,...,c_k) c₁,...,c_k FROM ...**.

10 | Summarizing Values: Aggregates

Aggregate functions (short: **aggregates**) reduce a *collection* of values to a *single* value (think summation, maximum).

- Simplest form: *collection* \equiv entire table:

```
SELECT agg1(e1) AS c1, ..., aggn(en) AS cn  
FROM ...
```

- Reduction of input rows: result table will have **one row**.
- Cannot mix aggregates with non-aggregate expression *e* in a **SELECT** clause:⁵ which value of *e* should we pick?

⁵ But see **GROUP BY** later on.

Aggregate Functions: Semantics

```
SELECT agg(e) AS c  -- e will typically refer to t
FROM   T AS t      -- range over entire table T
```

- Aggregate *agg* defined by triple $(\phi^{agg}, z^{agg}, \oplus^{agg})$:
 - ϕ^{agg} (*empty*): aggregate of the empty value collection
 - z^{agg} (*zero*): aggregate value initialiser
 - \oplus^{agg} (*merge*): add value to existing aggregate

```
a ←  $\phi^{agg}$            -- a will be aggregate value
for t in T             -- iterate over all rows of T
| x ← e(t)             -- value to be aggregated
| if x ≠ NULL          -- aggregates ignore NULL values (*)
| | if a =  $\phi^{agg}$       -- once we see first non-NULL value:
| | | a ←  $z^{agg}$         -- initialize aggregate
| | | a ←  $\oplus^{agg}(a, x)$  -- maintain running aggregate
```

Aggregate Functions: Semantics

Aggregate agg	ϕ^{agg}	z^{agg}	$\oplus^{agg}(a, x)$
COUNT	0	0	$a + 1$
SUM	NULL ⁶	0	$a + x$
AVG ⁷	NULL	$\langle 0, 0 \rangle$	$\langle a.1 + x, a.2 + 1 \rangle$
MAX	NULL	$-\infty$	$\max_2(a, x)$
MIN	NULL	$+\infty$	$\min_2(a, x)$
bool_and	NULL	true	$a \wedge x$
bool_or	NULL	false	$a \vee x$
\vdots	\vdots	\vdots	\vdots

- The special form **COUNT(*)** counts rows regardless of their fields' contents (NULL, in particular).

⁶ If you think “*this is wrong*,” we’re two already. Possible upside: sum differentiates between summation over an empty collection vs. a collection of all 0s.

⁷ Returns $a.1 / a.2$ as final aggregate value.

Aggregate Functions on Table T

```

SELECT COUNT(*)           AS "#rows",
       COUNT(t.d)         AS "#d",
       SUM(t.d)            AS "Σd",
       MAX(t.b)            AS "max(b)",
       bool_and(t.c)       AS "∀c",
       bool_or(t.d = 30)   AS "∃d=30"
FROM   T AS t
WHERE  p

```

#rows	#d	Σd	max(b)	∀c	∃d=30
5	4	100	'y'	false	true

$p \equiv \text{true}$

#rows	#d	Σd	max(b)	∀c	∃d=30
0	0	NULL	NULL	NULL	NULL

$p \equiv \text{false}$

Ordered Aggregates

- For most aggregates *agg*, merge \oplus^{agg} is commutative (and associative): row order does not matter.
- **Order-sensitive aggregates** admit a trailing `ORDER BY e_1, \dots, e_n` argument that defines row order:⁸

```
--          cast to text      separator string
--
SELECT string_agg(t.a :: text, ',' ORDER BY t.d) AS "all a"
FROM   T AS t
```

all a

'1,4,3,2,5'

⁸ $\oplus^{\text{string-agg}}$ essentially is `||` (string concatenation) which is not commutative.

Filtered and Unique Aggregates

```
SELECT agg(e) FILTER (WHERE p)
FROM ...
```

- **FILTER** clause alters aggregate semantics (see *):

$$\begin{array}{l} \vdots \\ x \leftarrow e(t) \\ \text{if } x \neq \text{NULL} \wedge p(x): \\ \vdots \end{array}$$

```
SELECT agg(DISTINCT e)
FROM ...
```

- Aggregates distinct (non-NULL) values of expression *e*.
(May use **ALL** to flag that duplicates are expected.)

11 : Forming Groups of Rows

Once **FROM** has generated row bindings, SQL clauses operate row-by-row. After **GROUP BY**: operate group-by-group:

SELECT e_1, \dots, e_m	-- 5
FROM ...	-- 1
WHERE ...	-- 2
GROUP BY g_1, \dots, g_n	-- 3
HAVING p	-- 4

- All rows that agree on all expressions g_i (the set of **grouping criteria**) form one **group**.
- \Rightarrow Steps 4 and 5 process groups (*not* individual rows). This affects expressions p and the e_j .

GROUP BY Partitions Rows

`SELECT ...`
`FROM ...`
`GROUP BY A_1`
`HAVING ...`

← evaluated once per **group** (not per row)

	A_1	A_2	...
	\vdots	\vdots	\vdots
the x_i group {	x_i	y_{i1}	\vdots
	x_i	y_{i2}	\vdots

the x_j group {	x_j	y_{j1}	\vdots
	x_j	y_{j2}	\vdots
	\vdots	\vdots	\vdots

Grouping **partitions** the row bindings:

- there are no empty groups
- each row belongs to exactly one group

GROUP BY Changes Field Types From τ To $\text{bag}(\tau)$ ⁹

👍	👎		**		👍
↓	↓		↓		↓
SELECT t.b, t.d FROM T AS t GROUP BY t.b			SELECT the(t.b) AS b, SUM(t.d) AS "Σd" FROM T AS t GROUP BY t.b		

- $t.d$ references current group of d values: violates 1NF!
⇒ After **GROUP BY**: **must** use aggregates on field values.
- $t.b$ references current group of b values **all of which are equal** in a group ⇒ SQL: using just $t.b$ is OK.
- (** May think of **hypothetical** aggregate $\text{the}(e)$ that picks one among equal e values.)

⁹ A view of **GROUP BY** that is due to Philip Wadler.

Aggregates are Evaluated Once Per Group

```

SELECT t.b                                AS "group",
       COUNT(*)                          AS size,
       SUM(t.d)                          AS "Σd",
       bool_and(t.a % 2 = 0)             AS "∀even(a)",
       string_agg(t.a :: text, ';' )    AS "all a"
FROM   T AS t
GROUP BY t.b;

```

group	size	Σd	∀even(a)	all a
'x'	3	40	false	'1;3;5'
'y'	2	60	true	'2;4'

- **HAVING** *p* acts like **WHERE** but *after* grouping: *p* = false discards groups (not rows).

Grouping Criteria

- The grouping criteria g_i form a set—order is irrelevant.
- Grouping on a **key** effectively puts each row in its own singleton group. (Typically a query smell 🤢.)
- Expressions e that are **functionally dependent** on the g_i are constant within a group (and thus can be used in `SELECT`).
 - If SQL does not know about the FD, explicitly add e to the set of g_i —this will *not* affect the grouping.

12 : Bag and Set Operations

Tables contain **bags of rows**. SQL provides the common family of binary **bag operations** (*no row order*):

q_1	UNION ALL	q_2	-- \cup^+ (bag union)
q_1	INTERSECT ALL	q_2	-- \cap^+ (bag intersection)
q_1	EXCEPT ALL	q_2	-- \setminus^+ (bag difference)

- Row types (width, field types) of the q_i must match.
- With **ALL**, row multiplicities are respected: if row r occurs n_i times in q_i , r will occur $\max(n_1 - n_2, 0)$ times in q_1 **EXCEPT ALL** q_2 (**INTERSECT ALL**: $\min(n_1, n_2)$).
 - Without **ALL**: obtain **set semantics** (no duplicates).

13 : Multi-Dimensional Data

- Relational representation of *measures* (*facts*) depending on multiple parameters (*dimensions*).
- Example: table `prehistoric` with *dimensions* `class`, `herbivore?`, `legs`, *fact* `species`:

<u>class</u>	<u>herbivore?</u>	<u>legs</u>	<u>species</u>
'mammalia'	true	2	'Megatherium'
'mammalia'	true	4	'Paraceratherium'
'mammalia'	false	2	NULL
'mammalia'	false	4	'Sabretooth'
'reptilia'	true	2	'Iguanodon'
'reptilia'	true	4	'Brachiosaurus'
'reptilia'	false	2	'Velociraptor'
'reptilia'	false	4	NULL

Table `prehistoric`

Multiple GROUP BYs: GROUPING SETS

- Analyze (here: group, then aggregate) table T along multiple dimensions \Rightarrow perform separate GROUP BYs on each relevant dimension:
- SQL syntactic sugar:


```
SELECT  $e_1, \dots, e_m$   
FROM  $T$   
GROUP BY GROUPING SETS ( $G_1, \dots, G_n$ )
```

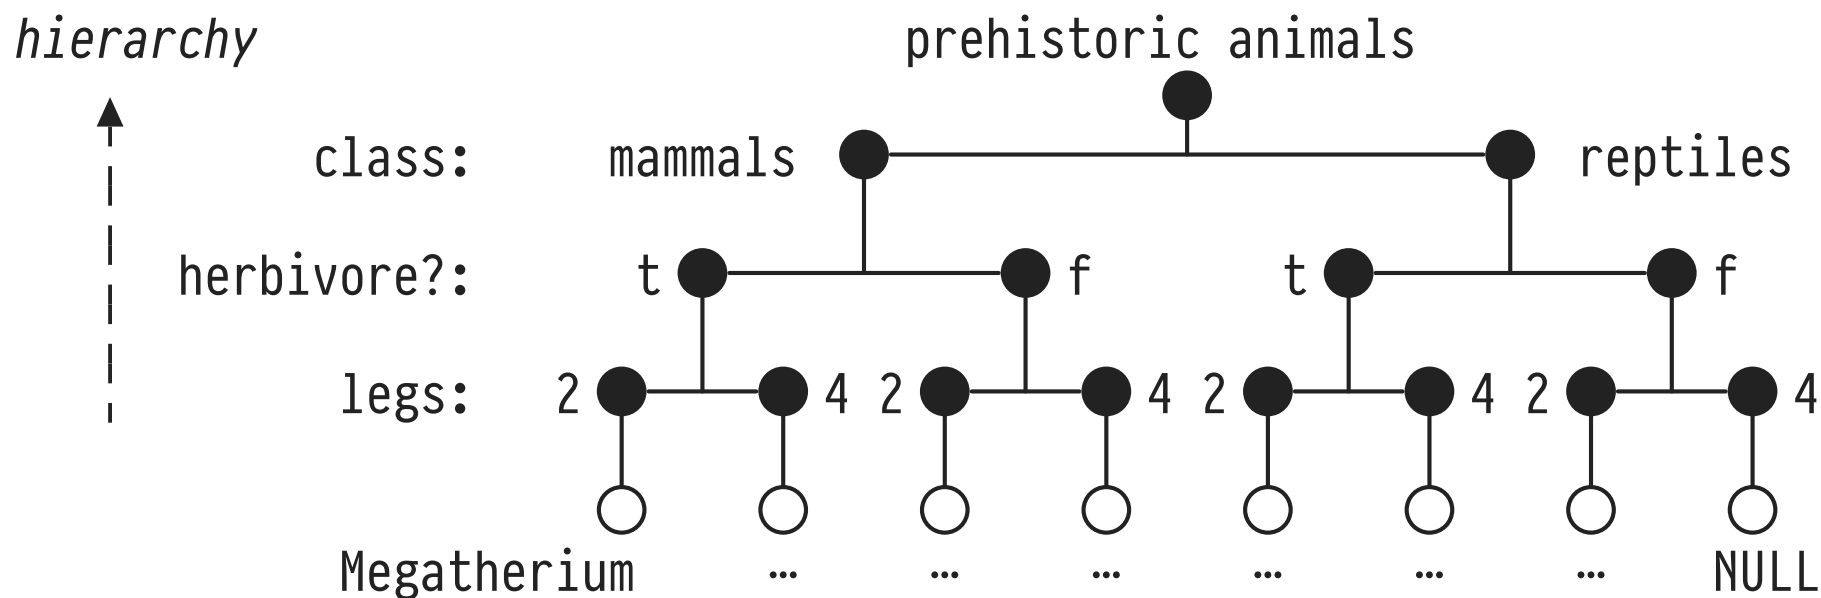
-- G_i : grouping criteria
-- sets in (...)

- Yields n GROUP BY queries q_i , glued together by UNION ALL. If non-aggregate $e_j \notin G_i$, $e_j \equiv \text{NULL}$ in q_i .

Hierarchical Dimensions: **ROLLUP**


- **Group along a path** from any node G_n up to the root:

ROLLUP $(G_1, \dots, G_n) \equiv$ GROUPING SETS $((G_1, \dots, G_{n-1}, G_n)$
 $, (G_1, \dots, G_{n-1}), \dots$
 $, (G_1),$
 $, ())$  ↓ *hierarchy*



Analyze All Dimension Combinations: CUBE

● \equiv slice for herbivore? = true



CUBE (G_1, \dots, G_n) \equiv GROUPING SETS	$\left(\begin{array}{c} (G_1, \dots, G_n) \\ \vdots \\ () \end{array} \right)$	} all 2^n subsets considered

14 : SQL Evaluation vs. Reading Order

```

SELECT DISTINCT ON (es 7) es 3, aggs 6
FROM   qs      1
WHERE  p        2
GROUP BY es     4
HAVING p        5

```

```

    UNION / EXCEPT / INTERSECT 8 } repeated 0 or more times,
    :                               } all evaluated before 9

```

```

ORDER BY es     9
OFFSET  n      10
LIMIT   n      10

```

- Reading order is: (7, 3, 6, 1 ⚠, 2, 4, 5, 8)⁺, 9, 10.

Query Nesting and (Non-)Readability

```
SELECT ...  
FROM (SELECT ...  
      FROM (SELECT ...  
            FROM ...  
            : ) AS descriptive  
      : ) AS ...  
:
```

- The more complex the query and the more useful the *descriptive* name becomes, the deeper it is buried. 🙄
- Query is a **syntactic monolith**. Tough to develop a query in stages/phases and assess the correctness of its parts.

15 | The **let...in** of SQL: **WITH** (Common Table Expressions)

Use **common table expressions (CTEs)** to bind table names *before* they are used, potentially multiple times:

WITH $T_1(c_{11}, \dots, c_{1, k_1})$ AS (q_1), \vdots $T_n(c_{n1}, \dots, c_{n, k_n})$ AS (q_n)	}	Query q_i may refer to tables T_1, \dots, T_{i-1}
q		q may refer to all tables T_i

- “Literate SQL”: Reading and writing order coincide.
- Think of **let** $T_1 = q_1, \dots, T_n = q_n$ **in** q in your favorite FP language. The T_i are undefined outside **WITH**.

SQL With **WITH** — Sample Uses


1. **Define queries in stages**, intermediate results in tables T_i . May use $q \equiv \text{TABLE } T_i^{10}$ to debug stage i .
2. **Bundle a query with test data:**

```
WITH
  prehistoric(class,"herbivore?",legs,species) AS (
    VALUES ('mammalia',true,2,'Megatherium'),
           ⋮
           ('reptilia',false,4,NULL)
  )
SELECT MAX(p.legs)
FROM   prehistoric AS p
```

¹⁰ Syntactic sugar for `SELECT t.* FROM T_i AS t.`

16 | 🔧 Use Case: **WITH** (Dinosaur Body Shapes)

Paleontology: **dinosaur body shape** (height/length ratio) and **form of locomotion** (using 2 or 4 legs) correlate:



- Use this correlation to infer bipedality (quadropedality) in incomplete dinosaur data sets:

<u>species</u>	height	length	legs
Gallimimus	2.4	5.5	?

Dinosaur Body Shapes

<u>species</u>	height	length	legs
Ceratosaurus	4.0	6.1	2
Deinonychus	1.5	2.7	2
Microvenator	0.8	1.2	2
Plateosaurus	2.1	7.9	2
Spinosaurus	2.4	12.2	2
Tyrannosaurus	7.0	15.2	2
Velociraptor	0.6	1.8	2
Apatosaurus	2.2	22.9	4
Brachiosaurus	7.6	30.5	4
Diplodocus	3.6	27.1	4
Supersaurus	10.0	30.5	4
Albertosaurus	4.6	9.1	NULL
Argentinosaurus	10.7	36.6	NULL
Compsognathus	0.6	0.9	NULL
Gallimimus	2.4	5.5	NULL
Mamenchisaurus	5.3	21.0	NULL
Oviraptor	0.9	1.5	NULL
Ultrasaurus	8.1	30.5	NULL

Table **dinosaurs**

Dinosaur Body Shapes

```
WITH
bodies(legs, shape) AS (
  SELECT d.legs, AVG(d.height / d.length) AS shape
  FROM   dinosaurs AS d
  WHERE  d.legs IS NOT NULL
  GROUP BY d.legs
)
:
```

<u>legs</u>	shape
2	0.447
4	0.201

Transient Table **bodies**

Dinosaur Body Shapes

- Query Plan:¹¹

0. Assume average body shapes in `bodies` are available
1. Iterate over all dinosaurs `d`:
 - If locomotion for `d` is known, output `d` as is
 - If locomotion for `d` is unknown:
 - Compute body shape for `d`
 - Find the `shape` entry `b` in `bodies` that matches `d`'s shape the closest
 - Use the locomotion (column `legs`) in `b` to complete `d`, output completed `d`

¹¹ In this course, *query plan* refers to a “plan of attack” for a query problem, not `EXPLAIN` output.